

Homework 4: Adversarial Attack on CIFAR-10

Implement White-Box PGD Attack and Black-Box Attack using PyTorch

```
In [ ]: import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
import numpy as np
from torch.autograd import Variable
import matplotlib.pyplot as plt
```

```
In [ ]: # 数据预处理
transform = transforms.Compose([
    transforms.ToTensor(),
])

# 加载 CIFAR-10 测试集
testset = torchvision.datasets.CIFAR10(
    root='./data', train=False, download=True, transform=transform)
testloader = torch.utils.data.DataLoader(
    testset, batch_size=1, shuffle=False, num_workers=2)
```

```
In [ ]: from pytorchcv.model_provider import get_model as ptcv_get_model
resnet20_cifar10 = ptcv_get_model("resnet20_cifar10", pretrained=True)
resnet56_cifar10 = ptcv_get_model("resnet56_cifar10", pretrained=True)

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
target_model = resnet20_cifar10.to(device) # 目标模型（攻击的对象）
surrogate_model = resnet56_cifar10.to(device) # 代理模型（黑盒攻击使用）
```

```
In [ ]: def test_clean_accuracy(model, test_loader, device):
    model.eval()
    correct = 0
    total = 0

    with torch.no_grad():
        for data, target in test_loader:
            data, target = data.to(device), target.to(device)
            output = model(data)
            pred = output.argmax(dim=1, keepdim=True)
            correct += pred.eq(target.view_as(pred)).sum().item()
            total += target.size(0)

    acc = 100. * correct / total
    print(f"[Clean Test] Accuracy: {correct}/{total} = {acc:.2f}%")
    return acc

# 测试目标模型原始精度
test_clean_accuracy(target_model, testloader, device)
```

[Clean Test] Accuracy: 8914/10000 = 89.14%

Out[]: 89.14

```
In [ ]: def pgd_attack(model, image, label, epsilon=8/255, alpha=2/255, iters=8):
        """
        PGD 对抗攻击
        :param model: 目标模型
        :param image: 原始图像 (1,C,H,W)
        :param label: 真实标签
        :param epsilon: 扰动上限 ( $L_\infty$ )
        :param alpha: 单步扰动强度
        :param iters: 迭代次数
        :return: 对抗样本
        """

        # 初始化对抗样本
        perturbed_image = image.clone().detach().requires_grad_(True)

        for _ in range(iters):
            # 前向传播
            output = model(perturbed_image)
            loss = nn.CrossEntropyLoss()(output, label)

            # 梯度清零并反向传播
            model.zero_grad()
            if perturbed_image.grad is not None:
                perturbed_image.grad.data.zero_()
            loss.backward()

            # 生成扰动
            data_grad = perturbed_image.grad.data
            sign_data_grad = data_grad.sign()

            # 更新对抗样本
            perturbed_image = perturbed_image + alpha * sign_data_grad

            # 投影到  $\epsilon$  邻域内
            perturbation = torch.clamp(
                perturbed_image - image,
                min=-epsilon,
                max=epsilon
            )
            perturbed_image = torch.clamp(image + perturbation, 0, 1).detach_()
            perturbed_image.requires_grad_(True)

        return perturbed_image.detach()

# 测试白盒攻击
def test_whitebox(model, testloader, epsilon=8/255):
    correct = 0
    total = 0

    for images, labels in testloader:
        images, labels = images.to(device), labels.to(device)

        # 生成对抗样本
        adv_images = pgd_attack(model, images, labels, epsilon)

        # 模型预测
        outputs = model(adv_images)
        _, predicted = torch.max(outputs.data, 1)
```

```

        total += labels.size(0)
        correct += (predicted == labels).sum().item()

    accuracy = 100 * correct / total
    print(f'White-box Attack Accuracy: {accuracy:.2f}% ( $\epsilon$ ={{epsilon}})')
    return accuracy

# 执行测试
test_whitebox(target_model, testloader)

```

White-box Attack Accuracy: 0.05% (ϵ =0.03137254901960784)

Out[]: 0.05

```

In [ ]: def blackbox_attack(target_model, surrogate_model, testloader, epsilon=8/255):
        correct = 0
        total = 0

        for images, labels in testloader:
            images, labels = images.to(device), labels.to(device)

            # 使用代理模型生成对抗样本
            adv_images = pgd_attack(surrogate_model, images, labels, epsilon)

            # 在目标模型上测试
            outputs = target_model(adv_images)
            _, predicted = torch.max(outputs.data, 1)

            total += labels.size(0)
            correct += (predicted == labels).sum().item()

        accuracy = 100 * correct / total
        print(f'Black-box Attack Accuracy: {accuracy:.2f}% ( $\epsilon$ ={{epsilon}})')
        return accuracy

# 执行黑盒攻击测试
blackbox_attack(target_model, surrogate_model, testloader)

```

Black-box Attack Accuracy: 59.43% (ϵ =0.03137254901960784)

Out[]: 59.43